

# lab\_music\_partial

November 15, 2021

## 1 Lab: Neural Networks for Music Classification

In addition to the concepts in the [MNIST neural network demo](#), in this lab, you will learn to: \*

- \* Load a file from a URL
- \* Extract simple features from audio samples for machine learning tasks such as speech recognition and classification
- \* Build a simple neural network for music classification using these features
- \* Use a callback to store the loss and accuracy history in the training process
- \* Optimize the learning rate of the neural network

To illustrate the basic concepts, we will look at a relatively simple music classification problem. Given a sample of music, we want to determine which instrument (e.g. trumpet, violin, piano) is playing. This dataset was generously supplied by [Prof. Juan Bello](#) at NYU Stenhardt and his former PhD student Eric Humphrey (now at Spotify). They have a complete website dedicated to deep learning methods in music informatics:

<http://marl.smusic.nyu.edu/wordpress/projects/feature-learning-deep-architectures/deep-learning-python-tutorial/>

You can also check out Juan's course.

### 1.1 Loading Tensorflow

Before starting this lab, you will need to install [Tensorflow](#). If you are using [Google colab](#), Tensorflow is already installed. Run the following command to ensure Tensorflow is installed.

```
conda activate tf2
```

```
[1]: import tensorflow as tf
```

Then, load the other packages.

```
[2]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
```

### 1.2 Audio Feature Extraction with Librosa

The key to audio classification is to extract the correct features. In addition to `keras`, we will need the `librosa` package. The `librosa` package in python has a rich set of methods extracting the features of audio samples commonly used in machine learning tasks such as speech recognition and sound classification.

Installation instructions and complete documentation for the package are given on the [librosa main page](#). On most systems, you should be able to simply use:

```
pip install librosa
```

For Unix, you may need to load some additional packages:

```
sudo apt-get install build-essential
sudo apt-get install libxext-dev python-qt4 qt4-dev-tools
pip install librosa
```

After you have installed the package, try to import it.

```
[3]: import librosa
import librosa.display
import librosa.feature
```

In this lab, we will use a set of music samples from the website:

<http://theremin.music.uiowa.edu>

This website has a great set of samples for audio processing. Look on the web for how to use the `requests.get` and `file.write` commands to load the file at the URL provided into your working directory.

You can play the audio sample by copying the file to your local machine and playing it on any media player. If you listen to it you will hear a soprano saxophone (with vibrato) playing four notes (C, C#, D, Eb).

```
[4]: import requests
fn = "SopSax.Vib.pp.C6Eb6.aiff"
url = "http://theremin.music.uiowa.edu/sound files/MIS/Woodwinds/
      ↪sopranosaxophone/"+fn

# TODO 1: Load the file from url and save it in a file under the name fn
```

```
[5]: r = requests.get(url)
with open(fn, 'wb') as f:
    f.write(r.content)
```

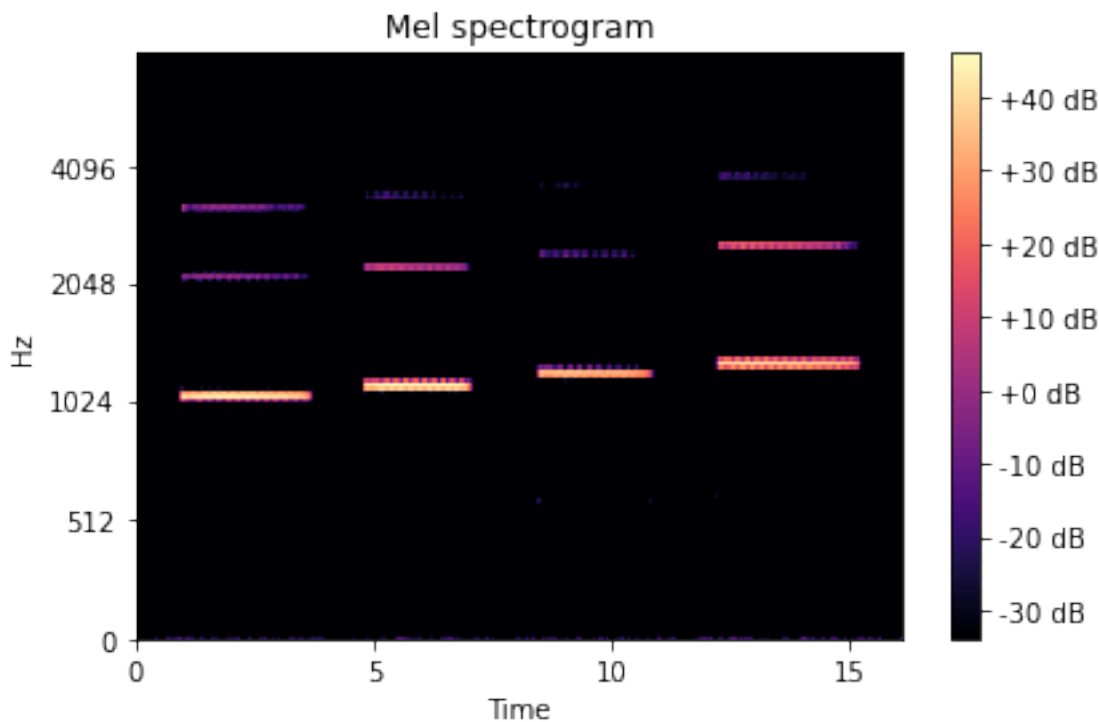
Next, use `librosa` command `librosa.load` to read the audio file with filename `fn` and get the samples `y` and sample rate `sr`.

```
[6]: # TODO 2
y, sr = librosa.load(fn)
```

Extracting features from audio files is an entire subject on its own right. A commonly used set of features are called the Mel Frequency Cepstral Coefficients (MFCCs). These are derived from the so-called mel spectrogram which is something like a regular spectrogram, but the power and frequency are represented in log scale, which more naturally aligns with human perceptual processing. You can run the code below to display the mel spectrogram from the audio sample.

You can easily see the four notes played in the audio track. You also see the ‘harmonics’ of each notes, which are other tones at integer multiples of the fundamental frequency of each note.

```
[7]: S = librosa.feature.melspectrogram(y=y, sr=sr, n_mels=128, fmax=8000)
librosa.display.specshow(librosa.amplitude_to_db(S),
                        y_axis='mel', fmax=8000, x_axis='time')
plt.colorbar(format='%+2.0f dB')
plt.title('Mel spectrogram')
plt.tight_layout()
```



### 1.3 Downloading the Data

Using the MFCC features described above, Eric Humphrey and Juan Bellow have created a complete data set that can be used for instrument classification. Essentially, they collected a number of data files from the website above. For each audio file, they segmented the track into notes and then extracted 120 MFCCs for each note. The goal is to recognize the instrument from the 120 MFCCs. The process of feature extraction is quite involved. So, we will just use their processed data provided at:

<https://github.com/marl/dl4mir-tutorial/blob/master/README.md>

Note the password. Load the four files into some directory, say `instrument_dataset`. Then, load them with the commands.

```
[8]: data_dir = 'instrument_dataset/'
Xtr = np.load(data_dir+'uiowa_train_data.npy')
ytr = np.load(data_dir+'uiowa_train_labels.npy')
Xts = np.load(data_dir+'uiowa_test_data.npy')
yts = np.load(data_dir+'uiowa_test_labels.npy')
```

Looking at the data files: \* What are the number of training and test samples? \* What is the number of features for each sample? \* How many classes (i.e. instruments) are there per class?

```
[9]: # TODO 3
num_train_sample, num_train_feature = np.shape(Xtr)
num_test_sample, num_test_feature = np.shape(Xts)
num_classes = len(np.unique(ytr))

print('num train sample = '+str(num_train_sample))
print('num test sample = '+str(num_test_sample))
print('num feature for each sample = '+str(num_train_feature))
print('num classes = '+str(num_classes))
```

```
num train sample = 66247
num test sample = 14904
num feature for each sample = 120
num classes = 10
```

Before continuing, you must scale the training and test data, `Xtr` and `Xts`. Compute the mean and std deviation of each feature in `Xtr` and create a new training data set, `Xtr_scale`, by subtracting the mean and dividing by the std deviation. Also compute a scaled test data set, `Xts_scale` using the mean and std deviation learned from the training data set.

```
[10]: # TODO 4: Scale the training and test matrices
# Xtr_scale = ...
# Xts_scale = ...
```

```
[11]: Xtr_mean = np.mean(Xtr, axis = 0)
Xtr_std = np.std(Xtr, axis = 0)
```

```
[12]: Xtr_scale = (Xtr - Xtr_mean)/Xtr_std
Xts_scale = (Xts - Xtr_mean)/Xtr_std
```

## 1.4 Building a Neural Network Classifier

Following the example in [MNIST neural network demo](#), clear the keras session. Then, create a neural network model with: \* `nh=256` hidden units \* `sigmoid` activation \* select the input and output shapes correctly \* print the model summary

```
[13]: from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import Dense, Activation
import tensorflow.keras.backend as K
```

```
[14]: # TODO 5 clear session
K.clear_session()
```

```
[15]: # TODO 6: construct the model
nin = num_train_feature # dimension of input data
nh = 256 # number of hidden units
nout = num_classes # number of outputs = 10 since there are 10 classes
model = Sequential()
model.add(Dense(units=nh, input_shape=(nin,), activation='sigmoid',
↳name='hidden'))
model.add(Dense(units=nout, activation='softmax', name='output'))
```

2021-11-15 09:46:03.666941: I tensorflow/core/platform/cpu\_feature\_guard.cc:145] This TensorFlow binary is optimized with Intel(R) MKL-DNN to use the following CPU instructions in performance critical operations: SSE4.1 SSE4.2 To enable them in non-MKL-DNN operations, rebuild TensorFlow with the appropriate compiler flags.  
2021-11-15 09:46:03.667454: I tensorflow/core/common\_runtime/process\_util.cc:115] Creating new thread pool with default inter op setting: 8. Tune using inter\_op\_parallelism\_threads for best performance.

```
[16]: # TODO 7: Print the model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
hidden (Dense)	(None, 256)	30976
output (Dense)	(None, 10)	2570

Total params: 33,546  
Trainable params: 33,546  
Non-trainable params: 0

Create an optimizer and compile the model. Select the appropriate loss function and metrics. For the optimizer, use the Adam optimizer with a learning rate of 0.001

```
[17]: # TODO 8
# opt = ...
# model.compile(...)
```

```
[18]: from tensorflow.keras import optimizers

opt = optimizers.Adam(lr=0.001)
model.compile(optimizer=opt,
```

```
loss='sparse_categorical_crossentropy',  
metrics=['accuracy'])
```

Fit the model for 10 epochs using the scaled data for both the training and validation. Use the `validation_data` option to pass the test data. Also, pass the callback class create above. Use a batch size of 100. Your final accuracy should be >99%.

```
[19]: # TODO 9  
hist = model.fit(Xtr_scale, ytr, epochs=10, batch_size=100,   
    ↪ validation_data=(Xts_scale, yts))
```

Train on 66247 samples, validate on 14904 samples

Epoch 1/10

66247/66247 [=====] - 3s 51us/sample - loss: 0.3514 -  
accuracy: 0.9051 - val\_loss: 0.2037 - val\_accuracy: 0.9344

Epoch 2/10

66247/66247 [=====] - 3s 41us/sample - loss: 0.0990 -  
accuracy: 0.9766 - val\_loss: 0.0865 - val\_accuracy: 0.9809

Epoch 3/10

66247/66247 [=====] - 3s 41us/sample - loss: 0.0585 -  
accuracy: 0.9858 - val\_loss: 0.0597 - val\_accuracy: 0.9862

Epoch 4/10

66247/66247 [=====] - 3s 40us/sample - loss: 0.0413 -  
accuracy: 0.9898 - val\_loss: 0.0484 - val\_accuracy: 0.9883

Epoch 5/10

66247/66247 [=====] - 3s 40us/sample - loss: 0.0314 -  
accuracy: 0.9920 - val\_loss: 0.0369 - val\_accuracy: 0.9895

Epoch 6/10

66247/66247 [=====] - 3s 40us/sample - loss: 0.0249 -  
accuracy: 0.9937 - val\_loss: 0.0316 - val\_accuracy: 0.9913

Epoch 7/10

66247/66247 [=====] - 3s 39us/sample - loss: 0.0205 -  
accuracy: 0.9945 - val\_loss: 0.0288 - val\_accuracy: 0.9917

Epoch 8/10

66247/66247 [=====] - 3s 38us/sample - loss: 0.0171 -  
accuracy: 0.9955 - val\_loss: 0.0324 - val\_accuracy: 0.9899

Epoch 9/10

66247/66247 [=====] - 3s 40us/sample - loss: 0.0146 -  
accuracy: 0.9964 - val\_loss: 0.0313 - val\_accuracy: 0.9908

Epoch 10/10


66247/66247 [=====] - 3s 40us/sample - loss: 0.0129 -  
accuracy: 0.9966 - val\_loss: 0.0320 - val\_accuracy: 0.9895

Plot the validation accuracy saved in `hist.history` dictionary. This gives one accuracy value per epoch. You should see that the validation accuracy saturates at a little higher than 99%. After that it “bounces around” due to the noise in the stochastic gradient descent.

```
[20]: # TODO 10
tr_accuracy = hist.history['accuracy']
val_accuracy = hist.history['val_accuracy']

plt.plot(tr_accuracy)
plt.plot(val_accuracy)
plt.grid()
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend(['training accuracy', 'validation accuracy'])
```

[20]: <matplotlib.legend.Legend at 0x7fe080208190>




Plot the loss values saved in the `hist.history` dictionary. You should see that the loss is steadily decreasing. Use the `semilogy` plot.

```
[21]: # TODO 11
tr_loss = hist.history['loss']
val_loss = hist.history['val_loss']

plt.semilogy(tr_loss)
plt.semilogy(val_loss)
plt.grid()
plt.xlabel('epochs')
plt.ylabel('loss')
```

```
plt.legend(['training loss', 'validation loss'])
```

```
[21]: <matplotlib.legend.Legend at 0x7fe0a88b4910>
```



## 1.5 Optimizing the Learning Rate

One challenge in training neural networks is the selection of the learning rate. Rerun the above code, trying four learning rates as shown in the vector `rates`. For each learning rate: \* clear the session \* construct the network \* select the optimizer. Use the Adam optimizer with the appropriate learning rate. \* train the model for 20 epochs \* save the accuracy and losses

```
[22]: rates = [0.01,0.001,0.0001]
batch_size = 100
loss_hist_tr = []
loss_hist_val = []
acc_hist_tr = []
acc_hist_val = []

# TODO 12
for lr in rates:
    K.clear_session()
    model = Sequential()
    model.add(Dense(units=nh, input_shape=(nin,), activation='sigmoid',
↪name='hidden'))
    model.add(Dense(units=nout, activation='softmax', name='output'))
```



```

opt = optimizers.Adam(lr=lr)
model.compile(optimizer=opt,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
hist = model.fit(Xtr_scale, ytr, epochs=20, batch_size=batch_size,
validation_data=(Xts_scale, yts))
loss_hist_tr.append(hist.history['loss'])
loss_hist_val.append(hist.history['val_loss'])
acc_hist_tr.append(hist.history['accuracy'])
acc_hist_val.append(hist.history['val_accuracy'])

```

Train on 66247 samples, validate on 14904 samples

Epoch 1/20

66247/66247 [=====] - 3s 47us/sample - loss: 0.1061 -  
accuracy: 0.9671 - val\_loss: 0.0403 - val\_accuracy: 0.9875

Epoch 2/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0299 -  
accuracy: 0.9904 - val\_loss: 0.0482 - val\_accuracy: 0.9811

Epoch 3/20

66247/66247 [=====] - 3s 41us/sample - loss: 0.0212 -  
accuracy: 0.9930 - val\_loss: 0.1433 - val\_accuracy: 0.9585

Epoch 4/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0188 -  
accuracy: 0.9939 - val\_loss: 0.1255 - val\_accuracy: 0.9701

Epoch 5/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0170 -  
accuracy: 0.9944 - val\_loss: 0.0294 - val\_accuracy: 0.9895

Epoch 6/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0134 -  
accuracy: 0.9955 - val\_loss: 0.0214 - val\_accuracy: 0.9923

Epoch 7/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0119 -  
accuracy: 0.9962 - val\_loss: 0.0348 - val\_accuracy: 0.9909

Epoch 8/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0135 -  
accuracy: 0.9955 - val\_loss: 0.0404 - val\_accuracy: 0.9896

Epoch 9/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0116 -  
accuracy: 0.9965 - val\_loss: 0.0532 - val\_accuracy: 0.9844

Epoch 10/20

66247/66247 [=====] - 3s 40us/sample - loss: 0.0155 -  
accuracy: 0.9951 - val\_loss: 0.0305 - val\_accuracy: 0.9921

Epoch 11/20

66247/66247 [=====] - 3s 39us/sample - loss: 0.0085 -  
accuracy: 0.9973 - val\_loss: 0.0248 - val\_accuracy: 0.9928

Epoch 12/20

66247/66247 [=====] - 3s 41us/sample - loss: 0.0083 -

accuracy: 0.9975 - val\_loss: 0.0625 - val\_accuracy: 0.9820  
 Epoch 13/20  
 66247/66247 [=====] - 3s 39us/sample - loss: 0.0084 -  
 accuracy: 0.9976 - val\_loss: 0.0991 - val\_accuracy: 0.9784  
 Epoch 14/20  
 66247/66247 [=====] - 3s 39us/sample - loss: 0.0093 -  
 accuracy: 0.9969 - val\_loss: 0.0508 - val\_accuracy: 0.9903  
 Epoch 15/20  
 66247/66247 [=====] - 3s 39us/sample - loss: 0.0096 -  
 accuracy: 0.9972 - val\_loss: 0.0482 - val\_accuracy: 0.9869  
 Epoch 16/20  
 66247/66247 [=====] - 3s 38us/sample - loss: 0.0077 -  
 accuracy: 0.9977 - val\_loss: 0.0460 - val\_accuracy: 0.9889  
 Epoch 17/20  
 66247/66247 [=====] - 3s 40us/sample - loss: 0.0091 -  
 accuracy: 0.9973 - val\_loss: 0.0376 - val\_accuracy: 0.9908  
 Epoch 18/20  
 66247/66247 [=====] - 3s 40us/sample - loss: 0.0059 -  
 accuracy: 0.9981 - val\_loss: 0.1059 - val\_accuracy: 0.9812  
 Epoch 19/20  
 66247/66247 [=====] - 3s 39us/sample - loss: 0.0109 -  
 accuracy: 0.9970 - val\_loss: 0.0522 - val\_accuracy: 0.9881  
 Epoch 20/20  
 66247/66247 [=====] - 3s 40us/sample - loss: 0.0077 -  
 accuracy: 0.9977 - val\_loss: 0.0460 - val\_accuracy: 0.9893  
 Train on 66247 samples, validate on 14904 samples  
 Epoch 1/20  
 66247/66247 [=====] - 3s 47us/sample - loss: 0.3634 -  
 accuracy: 0.9025 - val\_loss: 0.1992 - val\_accuracy: 0.9409  
 Epoch 2/20  
 66247/66247 [=====] - 3s 41us/sample - loss: 0.1030 -  
 accuracy: 0.9752 - val\_loss: 0.1087 - val\_accuracy: 0.9642  
 Epoch 3/20  
 66247/66247 [=====] - 3s 41us/sample - loss: 0.0600 -  
 accuracy: 0.9855 - val\_loss: 0.0568 - val\_accuracy: 0.9882  
 Epoch 4/20  
 66247/66247 [=====] - 4s 54us/sample - loss: 0.0425 -  
 accuracy: 0.9892 - val\_loss: 0.0594 - val\_accuracy: 0.9809  
 Epoch 5/20  
 66247/66247 [=====] - 3s 44us/sample - loss: 0.0320 -  
 accuracy: 0.9917 - val\_loss: 0.0394 - val\_accuracy: 0.9899  
 Epoch 6/20  
 66247/66247 [=====] - 3s 44us/sample - loss: 0.0253 -  
 accuracy: 0.9933 - val\_loss: 0.0323 - val\_accuracy: 0.9905  
 Epoch 7/20  
 66247/66247 [=====] - 3s 44us/sample - loss: 0.0207 -  
 accuracy: 0.9944 - val\_loss: 0.0297 - val\_accuracy: 0.9909  
 Epoch 8/20

66247/66247 [=====] - 3s 47us/sample - loss: 0.0175 - accuracy: 0.9955 - val\_loss: 0.0245 - val\_accuracy: 0.9932  
Epoch 9/20  
66247/66247 [=====] - 3s 46us/sample - loss: 0.0149 - accuracy: 0.9961 - val\_loss: 0.0247 - val\_accuracy: 0.9927  
Epoch 10/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0129 - accuracy: 0.9966 - val\_loss: 0.0259 - val\_accuracy: 0.9910  
Epoch 11/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0112 - accuracy: 0.9972 - val\_loss: 0.0257 - val\_accuracy: 0.9909  
Epoch 12/20  
66247/66247 [=====] - 3s 43us/sample - loss: 0.0105 - accuracy: 0.9974 - val\_loss: 0.0263 - val\_accuracy: 0.9908  
Epoch 13/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0092 - accuracy: 0.9976 - val\_loss: 0.0240 - val\_accuracy: 0.9913  
Epoch 14/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.0082 - accuracy: 0.9979 - val\_loss: 0.0184 - val\_accuracy: 0.9938  
Epoch 15/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.0073 - accuracy: 0.9982 - val\_loss: 0.0195 - val\_accuracy: 0.9926  
Epoch 16/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0069 - accuracy: 0.9982 - val\_loss: 0.0239 - val\_accuracy: 0.9910  
Epoch 17/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0064 - accuracy: 0.9984 - val\_loss: 0.0190 - val\_accuracy: 0.9927  
Epoch 18/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.0057 - accuracy: 0.9985 - val\_loss: 0.0219 - val\_accuracy: 0.9916  
Epoch 19/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0054 - accuracy: 0.9986 - val\_loss: 0.0179 - val\_accuracy: 0.9938  
Epoch 20/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0050 - accuracy: 0.9987 - val\_loss: 0.0233 - val\_accuracy: 0.9915  
Train on 66247 samples, validate on 14904 samples  
Epoch 1/20  
66247/66247 [=====] - 3s 51us/sample - loss: 1.1242 - accuracy: 0.6527 - val\_loss: 0.8697 - val\_accuracy: 0.6672  
Epoch 2/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.5620 - accuracy: 0.8433 - val\_loss: 0.5850 - val\_accuracy: 0.8118  
Epoch 3/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.3898 - accuracy: 0.9082 - val\_loss: 0.4579 - val\_accuracy: 0.8610

Epoch 4/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.3011 - accuracy: 0.9320 - val\_loss: 0.3581 - val\_accuracy: 0.8986

Epoch 5/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.2446 - accuracy: 0.9456 - val\_loss: 0.2884 - val\_accuracy: 0.9244

Epoch 6/20  
66247/66247 [=====] - 3s 44us/sample - loss: 0.2040 - accuracy: 0.9544 - val\_loss: 0.2496 - val\_accuracy: 0.9290

Epoch 7/20  
66247/66247 [=====] - 3s 49us/sample - loss: 0.1730 - accuracy: 0.9603 - val\_loss: 0.2143 - val\_accuracy: 0.9383

Epoch 8/20  
66247/66247 [=====] - 3s 42us/sample - loss: 0.1487 - accuracy: 0.9654 - val\_loss: 0.1786 - val\_accuracy: 0.9508

Epoch 9/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.1289 - accuracy: 0.9697 - val\_loss: 0.1517 - val\_accuracy: 0.9628

Epoch 10/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.1130 - accuracy: 0.9736 - val\_loss: 0.1331 - val\_accuracy: 0.9662

Epoch 11/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.1001 - accuracy: 0.9769 - val\_loss: 0.1275 - val\_accuracy: 0.9632

Epoch 12/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0894 - accuracy: 0.9796 - val\_loss: 0.1068 - val\_accuracy: 0.9731

Epoch 13/20  
66247/66247 [=====] - 3s 41us/sample - loss: 0.0806 - accuracy: 0.9813 - val\_loss: 0.0992 - val\_accuracy: 0.9733

Epoch 14/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0733 - accuracy: 0.9831 - val\_loss: 0.0907 - val\_accuracy: 0.9764

Epoch 15/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0672 - accuracy: 0.9841 - val\_loss: 0.0824 - val\_accuracy: 0.9802

Epoch 16/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0619 - accuracy: 0.9851 - val\_loss: 0.0774 - val\_accuracy: 0.9800

Epoch 17/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0574 - accuracy: 0.9863 - val\_loss: 0.0731 - val\_accuracy: 0.9803

Epoch 18/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0535 - accuracy: 0.9869 - val\_loss: 0.0675 - val\_accuracy: 0.9843

Epoch 19/20  
66247/66247 [=====] - 3s 40us/sample - loss: 0.0501 - accuracy: 0.9878 - val\_loss: 0.0644 - val\_accuracy: 0.9843

Epoch 20/20


66247/66247 [=====] - 3s 40us/sample - loss: 0.0471 - accuracy: 0.9883 - val\_loss: 0.0598 - val\_accuracy: 0.9859

Plot the loss function vs. the epoch number for all three learning rates on one graph. You should see that the lower learning rates are more stable, but converge slower.

```
[23]: # TODO 13
```

```
[24]: plt.semilogy(loss_hist_tr[0], label = 'lr = 0.01')
plt.semilogy(loss_hist_tr[1], label = 'r = 0.001')
plt.semilogy(loss_hist_tr[2], label = 'r = 0.0001')
plt.grid()
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend()
plt.title('training samples')
```


```
[24]: Text(0.5, 1.0, 'training samples')
```



```
[25]: plt.semilogy(loss_hist_val[0], label = 'lr = 0.01')
plt.semilogy(loss_hist_val[1], label = 'r = 0.001')
plt.semilogy(loss_hist_val[2], label = 'r = 0.0001')
plt.grid()
```

```
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend()
plt.title('test samples')
```

```
[25]: Text(0.5, 1.0, 'test samples')
```



```
[ ]:
```